

**INFOSYS SPRINGBOARD
VIRTUAL INTERNSHIP**



FIRE WEATHER INDEX PREDICTOR

SUBMITTED BY
SRINANDA C S

UNDER THE GUIDANCE OF
MENTOR PRAVEEN

PROJECT STATEMENT

This project focuses on developing a Machine Learning based Fire Weather Index (FWI) prediction system using meteorological parameters such as temperature, humidity, wind speed, rainfall, and fire-related indices. The dataset was preprocessed through data cleaning, handling missing values, feature scaling, and exploratory data analysis to ensure reliable model performance.

Multiple regression models were evaluated, and Ridge Regression was selected as the final model due to its better accuracy and ability to handle correlated features. The trained model was deployed using a Flask web application, enabling users to input weather conditions and obtain real-time FWI predictions along with corresponding fire risk levels.

EXPECTED OUTCOMES

- A predictive ML model trained using Ridge Regression to forecast FWI.
- A pre-processing pipeline using StandardScaler for normalization.
- A Flask-based web app where users can input environmental values and get FWI predictions.
- A system that can help forest departments, emergency planners, and climate researchers make data driven decisions.

MODULES TO BE IMPLEMENTED

- Data Collection
- Data Exploration (EDA) and Data Preprocessing
- Feature Engineering and Scaling
- Model Training using Ridge Regression
- Evaluation and Optimization
- Deployment via Flask App
- Presentation and Documentation

SYSTEM REQUIREMENTS

SOFTWARE REQUIREMENTS

- Operating System: Windows / Linux / macOS
- Programming Language: Python 3.11 (or above)
- Development Environment: Jupyter Notebook / VS Code / PyCharm
- Libraries & Frameworks:
 - NumPy – numerical computations
 - Pandas – data handling and preprocessing
 - Matplotlib & Seaborn – data visualization
 - Scikit-learn – machine learning model training and evaluation
 - Flask – web application development
 - Pickle – model and scaler serialization
- Web Browser: Google Chrome / Firefox / Microsoft Edge (for accessing the Flask app)

HARDWARE REQUIREMENTS

- Processor: Intel i3 / AMD Ryzen 3 or higher
- RAM: Minimum 4 GB (8 GB recommended)
- Storage: At least 5 GB free disk space
- System Type: 64-bit system
- Internet Connection: Required for dataset access and library installation

MILESTONE 1

Dataset Source : Kaggle

Importing Required Libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.model_selection import train_test_split,
GridSearchCV
from sklearn.linear_model import LinearRegression, Ridge,
Lasso
from sklearn.metrics import mean_absolute_error,
mean_squared_error, r2_score
import pickle
```

This code imports all the required Python libraries for data handling, visualization, preprocessing, model training, evaluation, and saving the trained model.

It prepares the environment for performing EDA, feature scaling, regression modeling (Ridge), performance evaluation, and model serialization.

Encoding Categorical Feature

```
df = pd.read_csv("FWI Dataset.csv")
if 'Region' in df.columns:
    df['Region'] = df['Region'].astype('category').cat.codes
```

Dataset is loaded using pandas.

The following information is displayed:

Entire dataset

Dataset structure (df.info())

Statistical summary (df.describe())

First & last 5 rows

Converts the Region column from strings to numerical category codes. Useful for machine learning algorithms that accept numeric inputs.

Selecting Numerical Features for Analysis

```
numeric_df = df.select_dtypes(include=['int64', 'float64'])
```

Extracts only numerical features for later analysis.

Initial Dataset Inspection

```
print(df.shape)  
print(df.columns)
```

To ensure the dataset structure is correct.

Identifying and Handling Missing Values

```
df.isnull().sum()  
df[df.isnull().any(axis=1)]  
df.columns = df.columns.str.strip()
```

Identifies missing values in each column.

Displays rows containing incomplete data.

Removes extra spaces in column names (common in raw CSV files).

Cleaning String Columns

```
for col in df.columns:  
    if df[col].dtype == 'object':  
        df[col] = df[col].astype(str).str.strip()
```

Removes unnecessary spaces in string values.

Ensures uniform data format.

Standardizing Text Data for Conversion

```
for col in df.columns:  
    if df[col].dtype == 'object':  
        df[col] = df[col].str.replace(" ", " ")  
    if df[col].str.contains(" ").any():  
        df[col] = df[col].str.split(" ").str[0]
```

This ensures numeric columns convert cleanly.

Converting Columns to Numeric Format

```
numeric_cols =  
['Temperature', 'RH', 'Ws', 'Rain', 'FFMC', 'DMC', 'DC', 'ISI', 'BUI',  
'FWI']  
for col in numeric_cols:  
    df[col] = pd.to_numeric(df[col], errors='coerce')
```

Converts corrupted strings into numeric format.
Non-convertible values become Nan.

Handling Missing Values in Categorical Columns

```
df['Region'] = df['Region'].fillna(df['Region'].mode()[0])
df['Classes'] = df['Classes'].fillna(df['Classes'].mode()[0])
```

Uses mode (most frequent value) for categorical features.
Prevents ML models from failing due to null values.

Label Encoding Categorical Columns

```
le_region = LabelEncoder()
df['Region_encoded'] = le_region.fit_transform(df['Region'])
le_class = LabelEncoder()
df['Classes_encoded'] = le_class.fit_transform(df['Classes'])
```

Convert string labels to numeric classes for ML model training.

Encoding All Remaining Categorical Columns

```
df_encoded = df.copy()
label_encoders = {}
for col in df_encoded.columns:
    if df_encoded[col].dtype == 'object':
        le = LabelEncoder()
        df_encoded[col] =
        le.fit_transform(df_encoded[col].astype(str))
        label_encoders[col] = le
```

This ensures all non-numeric features are usable in correlation analysis.



Correlation Heatmap

```
plt.figure(figsize=(10,8))
sns.heatmap(numeric_df.corr(), annot=True)
plt.show()
```

Used to understand:

Feature relationships


Which variables strongly influence FWI Multicollinearity issues



Histogram Plots

```
numeric_df.hist(figsize=(15, 12), bins=30)  
plt.show()
```


Shows distribution of each numeric feature:
Normal , Skewed, Outliers



Density (KDE) Plots

```
sns.kdeplot(numeric_df[col], fill=True)
```


These help understand:
Probability distribution, Spread of data, Detecting skewness



Boxplots for Outlier Detection

```
sns.boxplot(x=df[col])
```

Used to visually identify: Outliers, Data spread, Extreme values



Outlier Treatment using IQR

```

Q1 = numeric_df[col].quantile(0.25)
Q3 = numeric_df[col].quantile(0.75)
IQR = Q3 - Q1
df[col] = df[col].clip(lower, upper)
    
```

Purpose:

Removes extreme values

Prevents model distortion

Makes distributions more stable


Scatter Plots

```
sns.scatterplot(x=df['Temperature'], y=df['FWI'])
sns.scatterplot(x=df['Ws'], y=df['FWI'])
sns.scatterplot(x=df['RH'], y=df['FWI'])
```

Shows:

How individual features impact FWI

Linear / non-linear trends and Data clusters



To save the cleaned dataset.

```
df.to_csv("FWI_Cleaned.csv", index=False)
```

Milestone 1 Outcome Summary

- Fire Weather Index (FWI) dataset successfully loaded and inspected
- Dataset structure, data types, and summary statistics analyzed
- Missing values identified and handled using appropriate imputation techniques
- Categorical variables encoded for machine learning compatibility
- Data type casting applied to ensure numerical consistency
- Outliers detected using boxplots and distribution analysis

- Exploratory Data Analysis (EDA) performed using histograms, KDE plots, and scatterplots
- Correlation analysis conducted to understand feature relationships
- Cleaned and processed dataset saved as FWI Cleaned.csv for model development

MILESTONE 2

Data Loading and Preprocessing

```
df = pd.read_csv("FWI_Cleaned.csv")
```

The cleaned Fire Weather Index dataset prepared in Milestone 1 was loaded for model development.

```
df = df.dropna(subset=["FWI"])
```

Rows with missing target values were removed to avoid errors during model training.

Feature Selection and Target Definition

```
features =
["Temperature", "RH", "Ws", "Rain", "FFMC", "DMC", "DC", "ISI", "BUI"]
X = df[features]
y = df["FWI"]
```

Relevant meteorological and fire index features were selected based on correlation and domain relevance, with FWI chosen as the target variable.

Train-Test Split

```
train_test_split(X, y, test_size=0.2, random_state=42)
```

The dataset was split into 80% training and 20% testing data to ensure model generalization.

Feature Scaling

```
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
```

```
X_test_scaled = scaler.transform(X_test)
```

StandardScaler was applied to normalize feature values and ensure uniform scale across all input variables.

```
pickle.dump(scaler, f)
```

The trained scaler was saved as `scaler.pkl` to maintain consistency during model deployment.

Model Training

`LinearRegression()`, `Ridge()`, `Lasso()`

Linear, Ridge, and Lasso regression models were trained to compare performance and select the best model.

Model Performance Evaluation

`MAE`, `RMSE`, `R2 Score`

Model performance was evaluated using MAE, RMSE, and R² score to measure prediction accuracy and variance explanation.

Model Selection Based on Performance

```
results_df.sort_values(by="R2 Score")
```

A comparison table was created to objectively identify the best-performing regression model.

Hyperparameter Tuning

```
GridSearchCV(estimator=Ridge(),
param_grid={"alpha": [0.01, 0.1, 1, 10, 100]})
```

GridSearchCV was used to tune the alpha parameter to balance bias and variance.

Final Model Evaluation

```
best_ridge.predict(X_test_scaled)
```

The optimized Ridge model was evaluated on test data to confirm its performance.

```
pickle.dump(best_ridge, f)
```

The final trained Ridge Regression model was saved as *ridge.pkl* for deployment.

"Ridge Regression was selected..."

Ridge Regression was chosen due to its ability to handle multicollinearity among correlated weather features and its superior cross-validation performance.

Milestone 2 Outcome Summary

- Feature selection completed using correlation and domain knowledge
- Data normalized using StandardScaler
- Train–test split applied for generalization
- Multiple regression models evaluated
- Ridge Regression selected as the best model
- *scaler.pkl* and *ridge.pkl* successfully saved

MILESTONE 3

Loading Saved Model and Scaler

The previously trained and optimized Ridge Regression model and the StandardScaler are loaded using pickle to ensure consistency between training and evaluation.

```
with open("scaler.pkl", "rb") as f:  
    scaler = pickle.load(f)  
with open("ridge.pkl", "rb") as f:  
    ridge_model = pickle.load(f)
```

Dataset Loading for Evaluation

The cleaned dataset is loaded, and rows with missing target values are removed to maintain evaluation integrity.

```
df = pd.read_csv("FWI_Cleaned.csv")  
df = df.dropna(subset=["FWI"])
```

Dynamic Feature and Target Selection

Instead of hardcoding input features, all feature columns are selected dynamically by excluding the target variable.

```
target = "FWI"  
  
features = df.drop(columns=[target]).columns.tolist()  
  
X = df[features]  
  
y = df[target]
```

Train–Test Split

The dataset is split into training and testing subsets to evaluate the model on unseen data.

```
X_train, X_test, y_train, y_test = train_test_split(  
    X, y, test_size=0.2, random_state=42 )
```

Feature Scaling Using Saved Scaler

Only the test data is scaled using the previously saved scaler to avoid data leakage.

```
X_test_scaled = scaler.transform(X_test)
```

Model Prediction

Predictions are generated on the scaled test data using the finalized Ridge Regression model.

```
y_pred = ridge_model.predict(X_test_scaled)
```

Evaluation Metrics (MAE, RMSE, R² Score)

Multiple regression metrics are computed to evaluate prediction accuracy and error behavior.

```
mae = mean_absolute_error(y_test, y_pred)
rmse = np.sqrt(mean_squared_error(y_test, y_pred))
r2 = r2_score(y_test, y_pred)
```

Training vs Testing Performance (Accuracy)

Model performance is evaluated on both training and test data to check overfitting or underfitting.

```
train_r2 = ridge_model.score(
    scaler.transform(X_train), y_train
)
test_r2 = ridge_model.score(
    scaler.transform(X_test), y_test
)
```

Residual Analysis


Residuals are computed to analyze prediction errors.

```
residuals = y_test - y_pred
```

Residual Distribution Visualization

A histogram is used to analyze the distribution of residuals.

```
sns.histplot(residuals, kde=True)  
plt.axvline(0, color='red', linestyle='--')
```




Actual vs Predicted Visualization

A scatter plot compares predicted and actual FWI values using distinct colors.

Points close to the diagonal indicate strong predictive performance.

```
plt.scatter(y_test, y_pred, color="blue", alpha=0.6)  
plt.plot([y_test.min(), y_test.max()],  
[y_test.min(), y_test.max()],  
color="red", linestyle="--")
```




Residuals vs Predicted Values Plot

Residuals are plotted against predicted values to detect heteroscedasticity.

Random scatter confirms stable model behavior.

```
sns.scatterplot(x=y_pred, y=residuals, color="green")
plt.axhline(0, color="red", linestyle="--")
```



Best Alpha Parameter

The optimized alpha value selected during GridSearchCV is retrieved dynamically from the saved model.

This confirms that no manual hyperparameter tuning was performed during evaluation.

```
print("Final Ridge alpha used:", ridge_model.alpha)
```

All evaluation metrics are summarized in a structured format.

The Ridge Regression model demonstrates strong generalization on unseen data.

Residuals are centered around zero, and visualization confirms consistent predictive behavior.

Hyperparameter optimization effectively controlled multicollinearity among weather features.

Milestone 3 Outcome Summary

- Loaded the saved StandardScaler and optimized Ridge Regression model from Milestone 2 to ensure consistent preprocessing and evaluation.
- Used the cleaned FWI dataset and performed a train–test split to validate the model’s generalization on unseen data.
- Applied the saved scaler to both training and testing datasets to prevent data leakage.
- Generated predictions using the trained Ridge model on test data.
- Evaluated model performance using Mean Absolute Error (MAE) to measure average prediction error.
- Computed Root Mean Squared Error (RMSE) to penalize larger prediction errors and assess robustness.

- Calculated R² Score to measure how well the model explains variance in Fire Weather Index values.
- Computed training and testing R² scores to check for overfitting and confirm balanced learning.
- Performed residual analysis to examine prediction errors and verify that residuals are centered around zero.
- Visualized residual distribution to confirm unbiased error behavior.
- Plotted Actual vs Predicted FWI values to visually assess prediction accuracy.
- Analyzed Residuals vs Predicted values to ensure no systematic error patterns.
- Verified the final Ridge alpha value selected during hyperparameter tuning.
- Summarized evaluation metrics in a structured table for clear interpretation.
- Concluded that Ridge Regression provides stable, unbiased, and well-generalized predictions for Fire Weather Index estimation.

MILESTONE 4

FLASK WEB APPLICATION

app.py – Backend (Flask Application)

This file acts as the core backend of the web application. It loads the trained Ridge model and scaler, receives user inputs from the web form, preprocesses them, generates predictions, and sends results to the result page.

Importing Required Libraries

```
from flask import Flask, render_template, request
import numpy as np
```

```
import pickle
```

Flask is used to build the web application, NumPy handles numerical operations, and pickle loads the saved ML model and scaler.

Initializing Flask Application

```
app = Flask(__name__)
```

Creates the Flask application instance that handles routing and request processing.

Loading Trained Model and Scaler

```
with open("scaler.pkl", "rb") as f:  
    scaler = pickle.load(f)  
  
with open("ridge.pkl", "rb") as f:  
    model = pickle.load(f)
```

The saved StandardScaler and Ridge Regression model from Milestone 2 are loaded for real-time inference, ensuring consistency with training.

Defining Feature Order

```
FEATURES = [  
    "Temperature", "RH", "Ws", "Rain",  
    "FFMC", "DMC", "DC", "ISI", "BUI"  
]
```

The exact same feature names and order used in training are maintained to avoid mismatches during prediction.

Defining FWI Risk Categorization

```
def fwi_category(fwi):  
    if fwi <= 5:  
        return "Low"  
  
    elif fwi <= 11:
```

```

        return "Moderate"

    elif fwi <= 21:

        return "High"

    elif fwi <= 33:

        return "Very High"

    else:

        return "Extreme"

```

Converts numeric FWI prediction into a human-readable risk level, making the output meaningful for users.

Home Route

```
@app.route("/")
def index():

    return render_template("index.html")
```

Loads the main form page where users enter weather parameters.

Prediction Route

```
@app.route("/predict", methods=["POST"])

def predict():

    try:

        input_values = [float(request.form[f]) for f in
FEATURES]
```

Reads user inputs dynamically from the HTML form in the exact required order.

```
X = np.array([input_values])

X_scaled = scaler.transform(X)

prediction = model.predict(X_scaled) [0]
```

Transforms inputs using the saved scaler and feeds them into the trained Ridge model.

```
        return render_template(
            "home.html",
            prediction=round(float(prediction), 2),
            category=fwi_category(prediction)
        )
```

Sends prediction and risk category to the result page for display.

```
except Exception as e:
    return f"Prediction error: {e}"
```

Handles runtime errors instead of crashing the app.

Running the Flask Server

```
if __name__ == "__main__":
    app.run(debug=True)
```

Starts the local development server with debugging enabled.

index.html – Frontend (User Input Form)

This page collects weather inputs from the user and sends them to the Flask backend for prediction.

UI Styling

```
<style>
body {
    background-color: #121212;
    color: #ffffff;
    font-family: Arial;
}
```

```

form {
    width: 300px;
    margin: auto;
}

input, button {
    width: 100%;
    padding: 8px;
    margin: 6px 0;
    background: #1e1e1e;
    color: white;
    border: 1px solid #333;
}

button {
    background: #ff5722;
    border: none;
}

</style>

```

Implements a minimal dark theme for better aesthetics and readability.

Page Heading

```
<h2 style="text-align:center;">Fire Weather Index  
Prediction</h2>
```

Clearly states the purpose of the application to the user.

Input Form for Weather Features

```
<form method="POST" action="/predict">
```

Sends user inputs to the Flask /predict route using POST request.

```
<input type="number" step="any" name="Temperature"  
placeholder="Temperature (°C)" required>  
  
...  
  
<input type="number" step="any" name="BUI" placeholder="BUI"  
required>
```

Each input corresponds exactly to the trained ML model's features, ensuring compatibility.

```
<button type="submit">Predict FWI</button>
```

Triggers the prediction process when clicked.

home.html – Frontend (Prediction Result Page)

Displays the predicted FWI value along with a color-coded risk category.

Styling

```
<style>  
  
body {  
  
    background-color: #121212;  
  
    color: white;  
  
    font-family: Arial;  
  
    text-align: center;  
  
}  
  
.box {  
  
    margin-top: 100px;  
  
}  
  
.low { color: #4caf50; }  
  
.moderate { color: #ffeb3b; }  
  
.high { color: #ff9800; }
```

```
.veryhigh { color: #ff5722; }

.extreme { color: #f44336; }

</style>
```

Applies different colors to different risk levels for intuitive visualization.

Displaying Prediction

```
<h1>{{ prediction }}</h1>
```

Shows the numeric FWI value returned from the model.

Dynamic Risk Category Display

```
<h2 class="

  {% if category == 'Low' %}low
  {% elif category == 'Moderate' %}moderate
  {% elif category == 'High' %}high
  {% elif category == 'Very High' %}veryhigh
  {% else %}extreme{% endif %}
">

  Risk Level: {{ category }}
```

```
</h2>
```

Uses Flask templating to apply different colors based on risk level.

Navigation Back to Input Page

```
<a href="/" style="color:#ff5722;">Predict Again</a>
```

Allows users to easily return to the input page for another prediction.

← ⏪ 127.0.0.1:5000 Summarize ⚡ ⚡ ... Chat

Fire Weather Index Prediction

Temperature (°C)

Relative Humidity (%)

Wind Speed (km/h)

Rainfall (mm)

FFMC


DMC

DC


ISI

BUI

[Predict FWI](#)



WORKFLOW



CONCLUSION

In this project, a complete end-to-end Fire Weather Index (FWI) prediction system was successfully designed, implemented, evaluated, and deployed. Starting from raw data preprocessing and exploratory analysis, relevant meteorological and fire-related features were carefully selected and prepared for model training. Multiple regression models were evaluated, and Ridge Regression was chosen due to its superior performance and ability to handle multicollinearity among weather variables.

The model was optimized using hyperparameter tuning and validated using standard evaluation metrics such as MAE, RMSE, and R² score, ensuring reliable and generalized predictions. The trained model and scaler were serialized and seamlessly integrated into a Flask-based web application, enabling real-time FWI prediction through a user-friendly interface.

Overall, the system demonstrates an effective application of machine learning for environmental risk assessment, providing accurate fire risk predictions and clear risk categorization. This project not only validates the practical deployment of ML models but also highlights the importance of data preprocessing, model optimization, and scalable web integration for real-world decision-support systems.

