

Fire Weather Index Predictor

(A Machine Learning Model to Predict Fire Weather Index)

Infosys Springboard



Infosys SpringBoard Virtual Internship Program

Submitted by

S.MOUNIKA

Under the guidance of Mentor Praveen

PROBLEM STATEMENT:

We have developed a Fire Weather Index (FWI) Prediction System that will evaluate the potential risk of wildfires based on several key meteorological and fire behaviour factors. The FWI Prediction System will utilize temperature, humidity, wind speed, precipitation levels, as well as several important fire danger indices (FFMC, DMC, DC, ISI). The data has been thoroughly pre-processed, cleaned and explored to validate its quality and visually identify the most relevant trends in the data.

These advanced models allow us to model the relationships between multiple variables, both environmental and fire risk-related, for the creation of accurate and reliable Fire Weather Index Predictions (FWI). In addition to this input, we include both localized and regional-level information to help assess patterns of wildfires to improve the FWI Prediction System's applicability to local fire patterns. The ultimate objective of the FWI Prediction System is to assist both government agencies as well as decision-makers in finding fire locations before a major wildfire outbreak, taking preventative action and effectively allocating resources toward those areas at greatest risk.

EXPECTED OUTCOMES:

Deliverables are expected, but this list does not include everything. Some examples of deliverables are: a Ridge Regression model that can give accurate Fire Weather Index (FWI) predictions in different weather conditions; a data-preparation process that uses Standard Scaler to normalize the data so the model works better; a simple Flask web app where users can enter weather values and get FWI predictions instantly; and an analytical tool that helps forest departments, disaster management teams, and research groups plan for wildfires by using information from past data.

Project Modules :

- Gathering and organizing relevant dataset(s)
- Performing Exploratory Data Analysis and cleaning the data
- Engineering features and applying scaling techniques
- Building and training the model using Ridge Regression
- Assessing model performance and fine-tuning for improved accuracy
- Deploying the model through a user-friendly Flask web interface
- Creating final documentation and project presentation materials

Required Software :

- Python environment
- Python packages such as pandas, NumPy, matplotlib, seaborn, and others needed for analysis
- Flask for web application deployment

Milestone 1

Phase 1: Collecting Data :

Initially, the data to be used in making accurate predictions of the FWI was collected from different sources on the internet. The final choice was based on the environmental and fire danger parameter factors that were necessary or relevant to make an accurate prediction. The chosen dataset contained the following parameters: Temperature, Humidity, Wind Speed, Rainfall, FFMC, DMC, ISI, Regional Identification. All of these parameters became part of the dataset. Once the dataset was acquired, it was placed in a Pandas Data Frame. The first thing that was done was to check that the data was ready for analysis by reviewing its integrity.

The integrity review included assessing the types of data, identifying which components contained missing or inconsistent values, assessing how much memory the dataset consumed, generating statistical summaries for all variables, understanding the distribution of each feature within the dataset, confirming the number or number of rows and columns, and identifying and removing any duplicated rows. After confirming that the dataset had met all of the needed requirements for quality, the dataset was validated to ensure that it was accurate enough for the future phases of development and modelling.

1.loading the dataset

```
import pandas as pd df =  
pd.read_csv("fwi.xlsx.csv")  
print("loaded the dataset using pandas")
```

2. Verify Datatypes

Data Types of Each Column:

Feature	Data Type
Day	int64
Month	int64
Year	int64
Temperature	int64
RH	int64
Ws	int64
Rain	float64
FFMC	float64
DMC	float64
DC	object
ISI	float64
BUI	float64
FWI	object
Classes	object
Region	object

3. Basic Info of the Dataset

HEAD OF THE DATASET

Row	Day	Month	Year	Temperature	RH	Ws	Rain	FFMC	DMC	DC	ISI	BUI	FWI	Class	Region
0	1	6	2012	29	57	18	0.0	65.7	3.4	7.6	1.3	3.4	0.5	not fire	Bejaia
1	2	6	2012	29	61	13	1.3	64.4	4.1	7.6	1.0	3.9	0.4	not fire	Bejaia
2	3	6	2012	26	82	22	13.1	47.1	2.5	7.1	0.3	2.7	0.1	not fire	Bejaia
3	4	6	2012	25	89	13	2.5	28.6	1.3	6.9	0.0	1.7	0.0	not fire	Bejaia
4	5	6	2012	27	77	16	0.0	64.8	3.0	14.2	1.2	3.9	0.5	not fire	Bejaia

TAIL OF THE DATASET

Row	Day	Month	Year	Temperature	RH	Ws	Rain	FFMC	DMC	DC	ISI	BUI	FWI	Class	Region
239	26	9	2012	30	65	14	0.0	85.4	16.0	44.5	4.5	16.9	6.5	fire	Sidi-Bel Abbes
240	27	9	2012	28	87	15	4.4	41.1	6.5	8	0.1	6.2	0	not fire	Sidi-Bel Abbes
241	28	9	2012	27	87	29	0.5	45.9	3.5	7.9	0.4	3.4	0.2	not fire	Sidi-Bel Abbes
242	29	9	2012	24	54	18	0.1	79.7	4.3	15.2	1.7	5.1	0.7	not fire	Sidi-Bel Abbes
243	30	9	2012	24	64	15	0.2	67.3	3.8	16.5	1.2	4.8	0.5	not fire	Sidi-Bel Abbes

STATISTICAL SUMMARY OF THE DATASET

Statistic	Day	Month	Year	Temperature	RH	Ws	Rain	FFMC	DMC	ISI	BUI
Count	244	244	244	244	244	244	244	244	244	244	244
Mean	15.754 098	7.5000 00	2012 .0	32.172131	61.938 525	15.504 098	0.7606 56	77.887 705	14.673 361	4.7741 80	16.664 754
Std	8.8250 59	1.1129 61	0.0	3.633843	14.884 200	2.8101 78	1.9994 06	14.337 571	12.368 039	4.1753 18	14.204 824
Min	1.0	6.0	2012 .0	22.0	21.0	6.0	0.0	28.6	0.7	0.0	1.1
25%	8.0	7.0	2012 .0	30.0	52.0	14.0	0.0	72.075	5.8	1.4	6.0
50%	16.0	7.5	2012 .0	32.0	63.0	15.0	0.0	83.5	11.3	3.5	12.25
75%	23.0	8.0	2012 .0	35.0	73.25	17.0	0.5	88.3	20.75	7.3	22.525
Max	31.0	9.0	2012 .0	42.0	90.0	29.0	16.8	96.0	65.9	19.0	68.0

This table summarizes the principal attributes of the weather and fire danger derived from data in this data set. Important statistics for each attribute provided include: total number of values (count), average value (mean), amount which these values differ from one another (standard deviation), minimum & maximum value, & the percentiles (25%ile, 50%ile, & 75%ile). Percentiles define the distribution of these data.

In general, this summary will allow us to understand the extremes of the distribution, how the data is distributed around the mean, & what is typical in a fire season. Using this information will help

us to determine whether our data values appear reasonable before performing additional analysis/model building.

DUPLICATE VALUES :

```
print("\nChecking for Duplicate Rows:")  
print(df.duplicated().sum()) df=df.drop_duplicates()  
print("Shape after dropping duplicates:", df.shape)
```

Module 2: Data Exploration and Data Preprocessing

During the preprocessing and exploratory analysis phase, the dataset was carefully examined to ensure accuracy, completeness, and suitability for machine learning model development. The process involved handling missing values, detecting outliers, analyzing feature distributions, studying correlations, and encoding categorical attributes. The cleaned and transformed dataset was then prepared for use in subsequent modeling steps.

Handling Missing Values

A thorough check for missing or null values was performed across all columns.

The results were as follows:

Feature	Missing Values
Day	0
Month	0
Year	0
Temperature	0
RH	0
Ws	0
Rain	0
FFMC	0
DMC	0

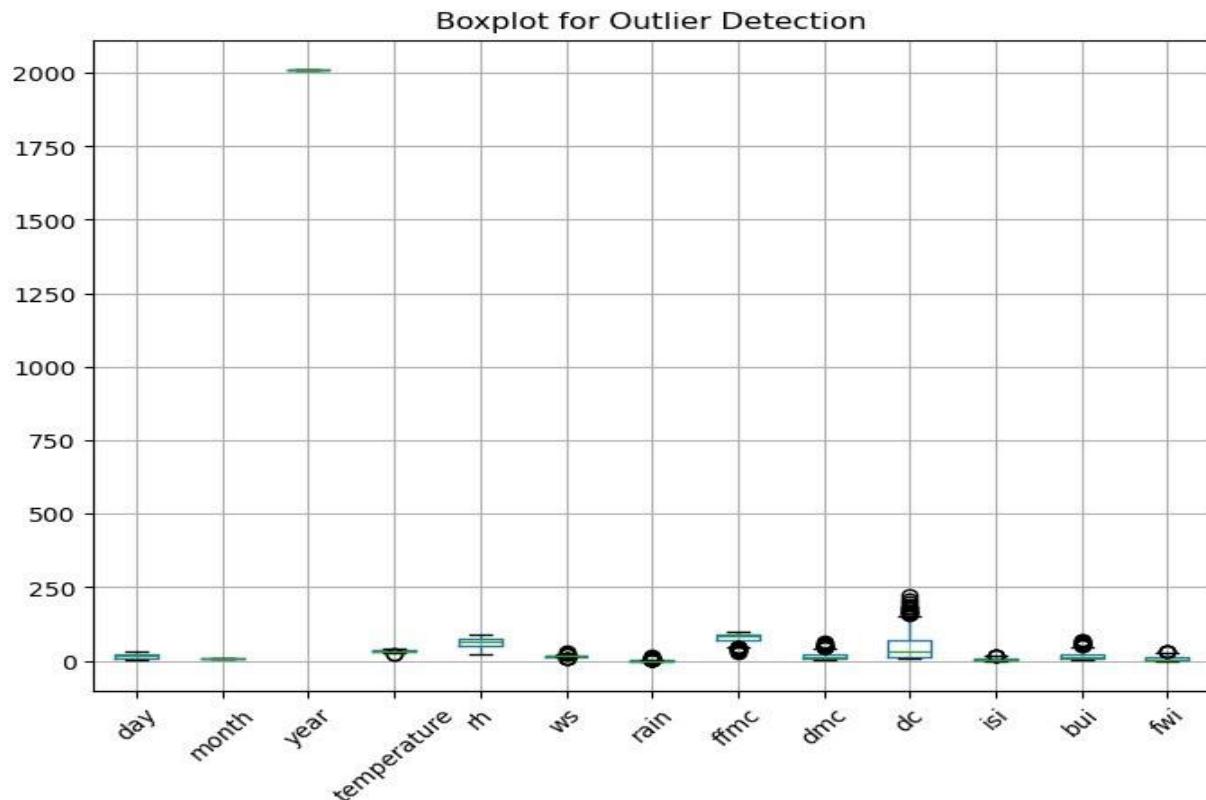
Feature	Missing Values
DC	1
ISI	0
BUI	0
FWI	1
Classes	1
Region	0

dtype: int64

Total rows with missing values: 1

Only one record contained a missing value in the *Classes* column. Identifying this early ensured data completeness and allowed for appropriate handling before further analysis.

Outlier Detection Using Boxplots and the IQR Method



(Fig1. outlier detection)

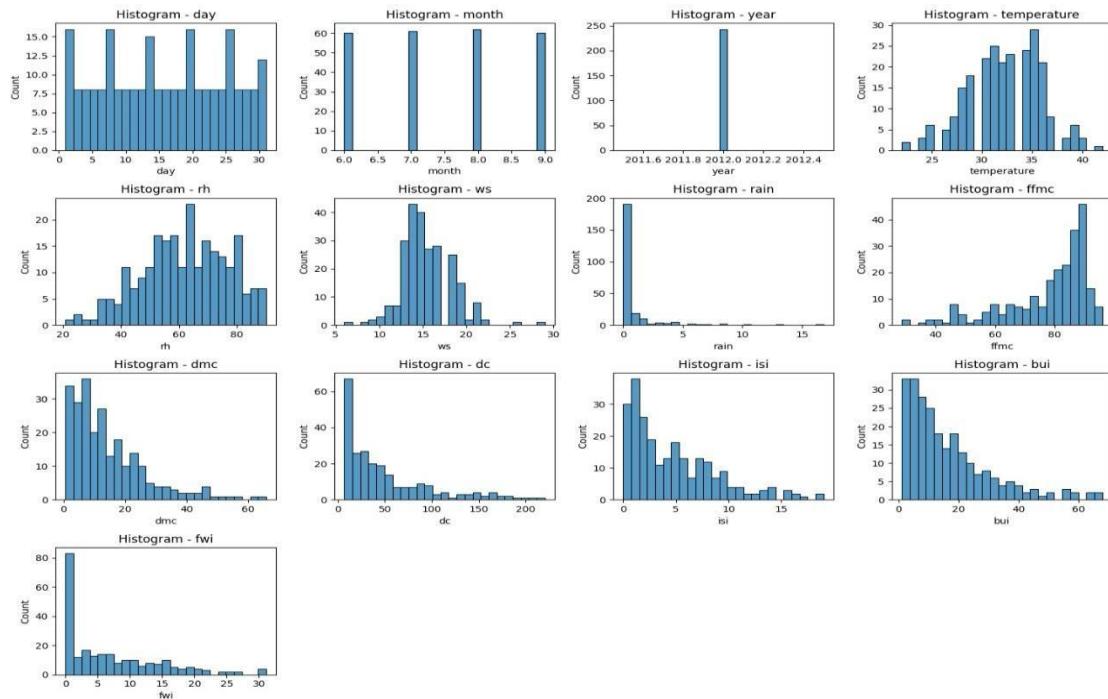
Outliers were examined using both visual (boxplots) and statistical (IQR-based) methods.

Key observations:


- **No outliers** were detected in: *day, month, year, RH, and DC*
- **Multiple outliers** were found in: *Temperature, Wind Speed, Rain, FFMC, DMC, ISI, and BUI*

These findings highlighted feature variability and supported decisions about cleaning and transformation.

3. Feature Distribution Analysis: Histograms and Density Plots



(FIG 2: Histograms)



(FIG.3 DENSITY PLOT)

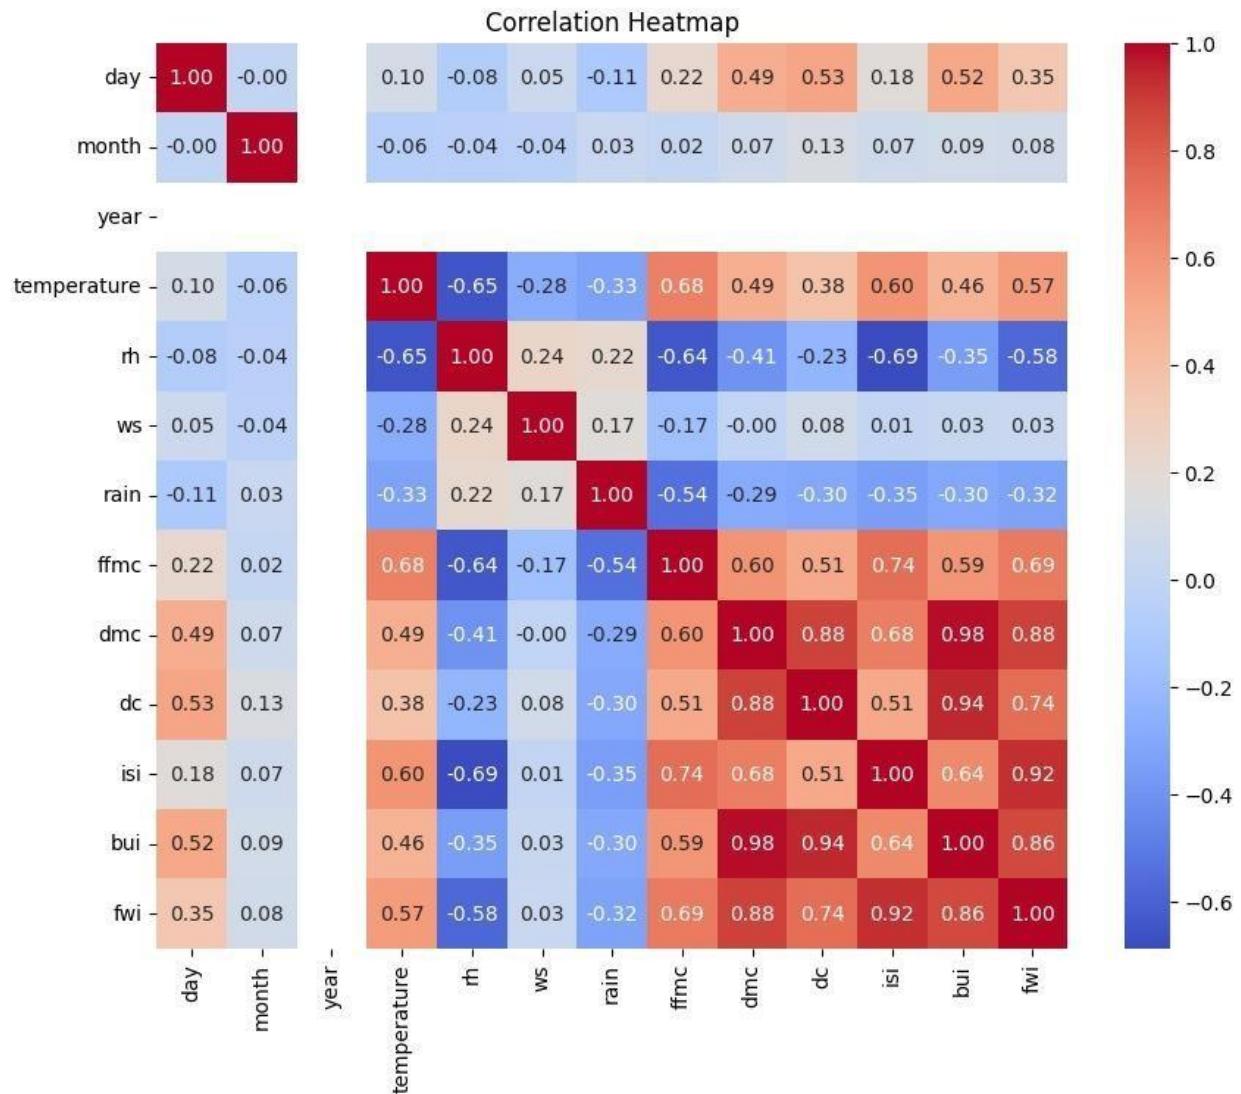
Histograms combined with KDE (Kernel Density Estimate) curves were generated for all numerical features to understand the overall distribution and behavior of the data. These visualizations help reveal important patterns that may influence the performance of the machine learning model.

Key Observations from the Graphs:

- **Temperature, Relative Humidity (RH), and Wind Speed** show smoother and more balanced distributions, indicating that these variables are spread fairly evenly across the dataset.
- **Rain, DMC, ISI, and BUI** exhibit strong right-skewness. This means most values are low, with only a few extreme high values present. Such skewness is common in environmental datasets where rainfall and fire indices vary significantly over time.

- Features like **FFMC** and **DC** display moderate variation with visible peaks, suggesting consistent seasonal or environmental patterns.
- The KDE curves helped identify whether each feature follows a normal distribution or deviates from it. Identifying non-normal distributions is crucial for selecting appropriate scaling methods and machine learning algorithms.

CORRELATION HEATMAP:



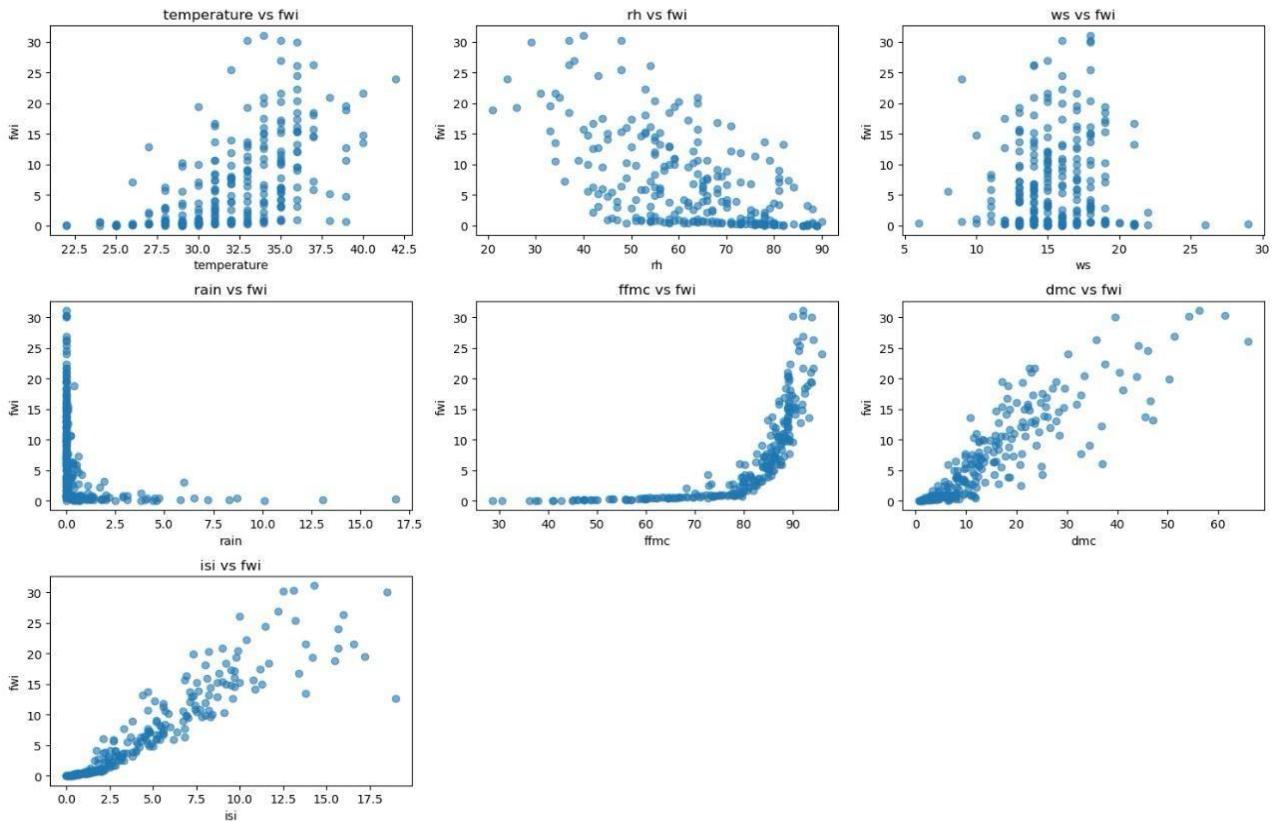
(FIG4:CORRELATION HEAT MAP)

A correlation heatmap was generated to study relationships among numerical features. Key findings:

- Strong positive correlations were observed among *BUI*, *DMC*, *FFMC*, *ISI*, and *Temperature*, suggesting shared patterns in fire-danger conditions.
- *Relative Humidity (RH)* showed strong negative correlations with several fire indices, reflecting its natural dampening effect on fire risk.

These insights helped identify the most influential predictors for the Fire Weather Index (FWI).

PAIR PLOTS FOR FEATURES VS FWI:



(FIG6:PAIR PLOTS)

Pair plots were generated to visually analyze the relationships between multiple numerical features in the dataset. These plots display scatterplots for every pair of variables along with individual feature distributions on the diagonal. This makes it easier to observe how different environmental and fire-danger features interact with each other.

- Variables such as **FFMC**, **DMC**, **ISI**, and **Temperature** show noticeable upward trends when paired with one another, indicating strong positive relationships linked to fire-prone conditions.
- **Relative Humidity (RH)** demonstrates an inverse pattern with several fire indices, reflecting the natural decrease in fire risk when humidity levels are higher.
- Some features, such as **Rain and Wind Speed**, show more scattered relationships, suggesting weaker correlations with the Fire Weather Index (FWI).
- The diagonal plots reveal each feature's distribution, helping identify skewness, spread, and potential outliers.

Encoding the Region Feature

To prepare categorical data for machine learning algorithms, the *Region* column was label encoded.

Region	Region Encoded
Bejaia	0

Region mapping: {'Bejaia': np.int64(0), 'Sidi-Bel Abbes': np.int64(1)} region

object

```
region_encoded    category dtype:  
object
```

SUMMARY OF THE TWO MODULES :

I focused on preparing the dataset for modeling by performing several data cleaning and analysis steps. I started by cleaning the column names to remove spaces, convert them to lowercase, and make them easier to work with. After loading the dataset, I checked for missing values across all rows and columns, and confirmed that there were no null values present.

Next, I converted important numeric columns like DC and FWI into float type to ensure they were ready for mathematical operations. The region column was then label-encoded so it could be used in machine learning models. Since the classes column was not required for the regression task, I removed it from the dataset.

After cleaning, I generated boxplots to visually detect outliers in the numeric features and used scatterplots to study the relationship between each feature and the FWI value. I also created a correlation heatmap to understand how different features are related to one another

Milestone 2

Module 3: Feature Engineering and scaling:

The main objective of Module 3 is to prepare the dataset for machine learning model training. After cleaning the data in the previous module, important features were selected and their scale was normalized. Based on correlation analysis and basic domain understanding of the Fire Weather Index (FWI), important features such as Temperature, Relative Humidity (RH), Wind Speed, Rain, FFMC, DMC, DC, ISI, and BUI were selected as input variables, while FWI was taken as the target variable.

Unnecessary date-related columns such as day, month, and year were removed because they do not directly affect fire behaviour. The dataset was then divided into input features (X) and the target variable (y). To make sure that all numerical features are on the same scale and contribute equally during model training, feature scaling was performed using Standard Scaler.

Finally, the fitted scaler was saved as a .pkl file so that the same preprocessing steps can be applied during testing and future deployment. This helps in maintaining consistency while handling real-time inputs.

1. Loading the cleaned dataset into pandas:

```
df=pd.read_csv("cleaned_FWI_dataset.csv")
print("Reloaded cleaned dataset:")
```

2. Feature Selection using Correlation Analysis:

Correlation analysis was performed and features with absolute correlation greater than or equal to 0.3 with the Fire Weather Index (FWI) were selected as input features.

```
Index(['fwi', 'isi', 'dmc', 'bui', 'dc', 'ffmc', 'rh', 'temperature',
'rain'], dtype='object')
```

3. The dataset is split into training and testing sets using a train test split.

```
x_train, x_test, y_train, y_test = train_test_split(  
    X, y, test_size=0.2, random_state=42)
```

The dataset is split into training and testing sets to evaluate model performance on unseen data.

4. Feature Scaling using Standard Scaler:

Standard Scaler is used to normalize the input features so that all numerical values are on the same scale. It transforms the data such that each feature has a mean of 0 and a standard deviation of 1. This step is important for linear and regularized regression models, as it helps improve model performance and ensures that no feature dominates due to larger values.

Shape of scaled training data: (194, 8)

Statistical summary of scaled training data:

Statistic	isi	dmc	bui	dc	ffmc
Count	1.940000e+02	1.940000e+02	1.940000e+02	1.940000e+02	1.940000e+02
Mean	-4.921607e-17	-7.325183e-17	-9.156479e-17	9.614302e-17	-1.327689e-16
Std	1.002587e+00	1.002587e+00	1.002587e+00	1.002587e+00	1.002587e+00
Min	-1.143565e+00	-1.095782e+00	-1.060395e+00	-8.894062e-01	-3.514355e+00
25%	-8.138422e-01	-7.204398e-01	-7.529531e-01	-8.216428e-01	-4.968853e-01
50%	-2.990111e-01	-3.067969e-01	-3.391536e-01	-3.633611e-01	3.835060e-01
75%	6.091740e-01	4.841035e-01	4.352665e-01	5.150539e-01	7.669022e-01
Max	3.252745e+00	3.883252e+00	3.366762e+00	3.397259e+00	1.270997e+00

The trained Standard Scaler was saved as a .pkl file. This ensures that the same scaling process can be applied to test data and future inputs during model deployment, maintaining consistency in preprocessing.

Module 4: Model Training using Ridge Regression:-

Multiple supervised regression models were trained and evaluated to develop an effective FWI prediction system. Scaled data was used for training, and GridSearchCV was applied for hyperparameter tuning where necessary. The comparative approach helped select a model that generalizes well and minimizes prediction error.

1.Linear Regression:

- A supervised learning algorithm used to predict FWI.
- Trained using scaled input features
- Used as a baseline model for comparison

```
lr = LinearRegression()  
lr.fit(X_train_scaled, y_train)
```

2.Ridge Regression:

- A regularized version of Linear Regression.
- Helps reduce overfitting.
- Hyperparameters were tuned using GridSearchCV
- Best model saved as ridge.pkl for deployment

```
ridge = Ridge()  
ridge.fit(X_train_scaled, y_train)
```

3.Lasso Regression:

- Performs feature selection by shrinking coefficients.
- Helps handle multicollinearity.
- Trained using scaled data.

```
lasso = Lasso(alpha=0.01)

lasso.fit(X_train_scaled, y_train)
```

4.ElasticNet Regression:

- Combines the advantages of both Ridge and Lasso regression.
- Works well when input features are correlated.
- Trained using scaled input features.

```
enet = ElasticNet(alpha=0.01, l1_ratio=0.5)

enet.fit(X_train_scaled, y_train)
```

5.Decision Tree Regressor:

- A tree-based regression model used for prediction.
- Captures non-linear relationships in the data.
- Does not require feature scaling.

```
t = DecisionTreeRegressor(random_state=42)

dt.fit(X_train, y_train)
```

Training and testing accuracy were evaluated using R² score:

Model Accuracy Results:

Linear Regression:

Train: 0.972975778450515

Test : 0.9899092711837978

Ridge Regression:

Train: 0.9728269417813818

Test : 0.9828463110256618

Lasso Regression:

Train: 0.9729294016765321

Test : 0.9880770054667244

ElasticNet Regression:

Train: 0.9728084514553315

Test : 0.9828083385032822

Decision Tree Regressor:

Train: 1.0

Test : 0.9330959844728333

Ridge Regression was further improved using hyperparameter tuning:

```
ridge = Ridge()  
  
# Alpha values to try  
  
params = {  
    'alpha': [0.01, 0.1, 1, 10, 100]  
}  
  
# Apply GridSearchCV  
  
ridge_gs = GridSearchCV(  
    ridge,
```

```
    params,  
    cv=5,  
    scoring='r2'  
)  
  
# Fit on training data  
ridge_gs.fit(X_train_scaled, y_train)  
  
  
# Best Ridge model  
best_ridge = ridge_gs.best_estimator_  
  
  
print("Best alpha value:", ridge_gs.best_params_)
```

All trained regression models were compared using testing accuracy (R^2 score):

Model Comparison (Testing R^2):

Linear Regression : 0.9899092711837978


Lasso Regression : 0.9880770054667244

ElasticNet Regression : 0.9828083385032822

Decision Tree : 0.9330959844728333

Best Model (Tuned Ridge Regression): 0.9670235224654835

An Actual vs Predicted plot was used to visualize model performance:



SUMMARY OF THE TWO MODULES:

Module 3 focused on preparing the cleaned dataset for machine learning. Important features were selected using correlation analysis, the target variable was separated, and the data was split into training and testing sets. Feature scaling was applied using Standard Scaler, and the scaler was saved to ensure consistent preprocessing for future use.

Module 4 focused on building and evaluating machine learning models. Multiple regression models were trained, and their performance was evaluated using training and testing accuracy. Ridge Regression was

further tuned using GridSearchCV to improve performance. All models were compared to select the best model, performance was visualized using a plot chart, and the trained models were saved for future use.

Milestone 3

Module5: Evaluation and Optimization

In this module, the trained Ridge Regression model is evaluated to analyze its prediction performance. Predictions are generated on the testing dataset using the trained model.

The model performance is evaluated using standard regression metrics such as **Mean Absolute Error (MAE)**, **Root Mean Squared Error (RMSE)**, and **R² Score**. These metrics help in understanding the prediction error and how well the model explains the variance in the Fire Weather Index.

To further analyze the model performance, an **Actual vs Predicted visualization** is plotted. This plot visually compares the predicted values with the actual values and helps in assessing the accuracy of the model. The close alignment between actual and predicted values indicates good model performance.

Based on the evaluation results, the Ridge Regression model demonstrates reliable performance and is suitable for Fire Weather Index prediction.

Evaluated the models using mean absolute error (MAE):

The models were evaluated using Mean Absolute Error (MAE) to measure the average difference between the actual and predicted values. MAE provides a clear understanding of how much the model's predictions deviate from the true values. A lower MAE indicates better prediction accuracy, making it useful for comparing the performance of different regression models.

Model	Train MAE	Test MAE
Linear Regression	6.818935e-01	0.412130
Ridge Regression	6.883442e-01	0.463768
Lasso Regression	6.793818e-01	0.425542

Model	Train MAE	Test MAE
ElasticNet Regression	6.884746e-01	0.463063
Decision Tree Regression	7.868849e-18	0.879592

RMSE is used to evaluate the prediction error of the models:

RMSE is used to evaluate the prediction error of the models by measuring the square root of the average squared difference between the actual and predicted values. It gives more weight to larger errors, which helps in understanding how well the model handles significant prediction mistakes. A lower RMSE value indicates better model performance.

Index	Model	Train RMSE	Test RMSE
0	Linear Regression	1.280460e+00	0.553023
1	Ridge Regression	1.283981e+00	0.721042
2	Lasso Regression	1.281558e+00	0.601138
3	ElasticNet Regression	1.284418e+00	0.721839
4	Decision Tree Regression	2.621033e-17	1.444695

Evaluated the models performance using the R² score:

The models performance was evaluated using the R² score, which measures how well the model explains the variance in the target variable. An R² value closer to 1 indicates that the model provides a better fit to the data and makes accurate predictions.

Model	Train R²	Test R²
Linear Regression	0.972976	0.989909
Ridge Regression	0.972827	0.982846
Lasso Regression	0.972929	0.988077
Elastic Net Regression	0.972808	0.982808
Decision Tree Regression	1.00000	0.931137

Hyperparameter Tuning and Model Saving:

In this step, Ridge Regression is optimized by tuning the alpha parameter using GridSearchCV. The best model is selected based on error metrics, evaluated using MAE, RMSE, and R² score, and then saved using joblib for future use.

Best Alpha: {'alpha': 0.1}

MAE: 0.42659747618715266

RMSE: 0.5978891362204466

R² Score: 0.9882055415087834

Tuned Ridge model saved successfully!

MAE Calculation:

$$\text{MAE} = (1 / n) \times \sum |y_i - \hat{y}_i|$$

Root Mean Squared Error (RMSE):

$$\text{RMSE} = \sqrt{[(1 / n) \times \sum (y_i - \hat{y}_i)^2]}$$

R² Score:

$$R^2 = 1 - [\sum (y_i - \hat{y}_i)^2 / \sum (y_i - \bar{y})^2]$$

Summary of module 5:

In Module 5, the trained regression models were evaluated to analyze their prediction performance. Model predictions were generated on the test dataset and evaluated using standard metrics such as Mean Absolute Error (MAE), Root Mean Squared Error (RMSE), and R² score. These metrics helped in understanding the accuracy and reliability of the models.

An Actual vs Predicted visualization was used to visually assess the model's prediction performance. Additionally, the performance of different regression models was compared using R² scores. Based on the evaluation results, Ridge Regression showed stable and consistent performance and was selected as the final model for Fire Weather Index prediction.